

MachineLearning _Clustering

1

-Rathiga Ramesh

Full Stack Developer, AI Enthusiast

Rathiga_ML_Cluster

Affinity Propagation

How it works:

- Uses "message passing" between data points to identify **exemplars** (most representative points).
- Does **not require pre-specifying** the number of clusters.

Strengths:

- Works well with **small to medium-sized datasets**.
- Finds **natural clusters** without forcing a fixed structure.


Weaknesses:

- **High computational cost** ($O(N^2)$ memory/time complexity).
- Sensitive to **damping factor** and **preference parameter**.

Best for:

- Small datasets where the number of clusters is unknown (e.g., gene expression, image segmentation).

```
from sklearn.cluster import AffinityPropagation
aff = AffinityPropagation(random_state=5)
y_pred = aff.fit_predict(X)
```



AgglomerativeClustering (Hierarchical Clustering)

How it works:

- **Bottom-up approach:** Starts with each point as a cluster, then **merges closest pairs** iteratively.
- Uses **linkage criteria** (ward, complete, average, single).

Strengths:

- Produces a **dendrogram** for multi-level clustering analysis.
- Flexible with **different distance metrics**.

Weaknesses:


- **Not scalable** for large datasets ($O(N^3)$ time complexity).
- Once merged, clusters cannot be split.

Best for:

- Medium-sized datasets where hierarchy matters (e.g., taxonomy, document clustering).

Rathiga_ML_Cluster

```
from sklearn.cluster import AgglomerativeClustering
agg = AgglomerativeClustering(n_clusters=3)
y_pred = agg.fit_predict(X)
```



K-Means

How it works:

- Partitions data into **K clusters** by minimizing **inertia** (within-cluster variance).
- Uses **iterative centroid updates**.

Strengths:

- Fast and scalable** ($O(N \cdot K)$ per iteration).
- Works well with **spherical clusters**.


Weaknesses:

- Sensitive to initialization** (k-means++ helps).
- Struggles with **non-convex clusters**.

Best for:

- Large datasets with clear separation (e.g., market segmentation, image compression).

```
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=3, random_state=5)
y_pred = kmeans.fit_predict(X)
```



Mean Shift

How it works:

- ▶ **Kernel density estimation** to find cluster modes.
- ▶ **Shifts points** towards high-density regions.

Strengths:

- ▶ **No predefined cluster count** needed.
- ▶ Robust to **outliers**.

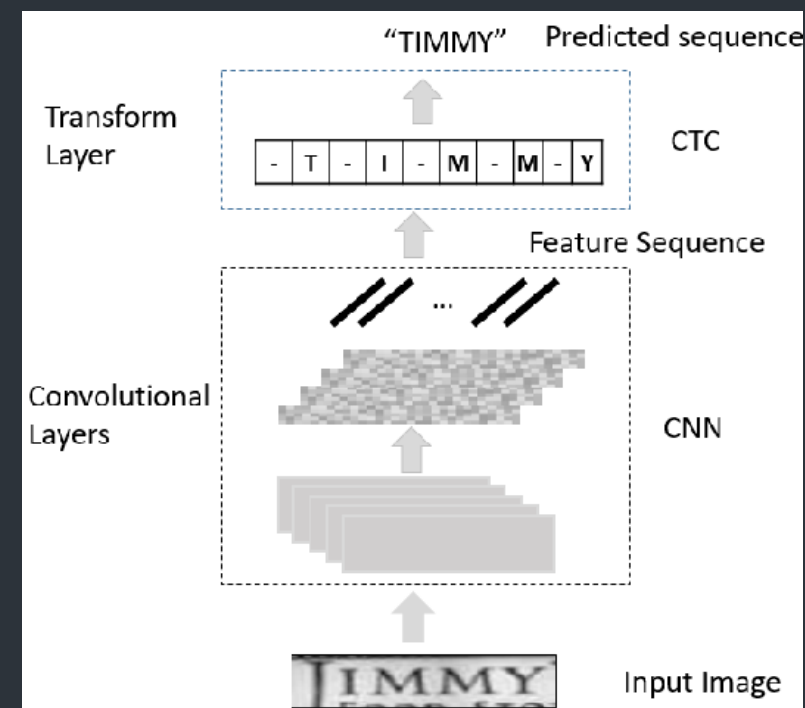
Weaknesses:

- ▶ **Computationally expensive** ($O(N^2)$).
- ▶ **Bandwidth selection** is critical.

Best for:

- ▶ Image segmentation, object tracking.

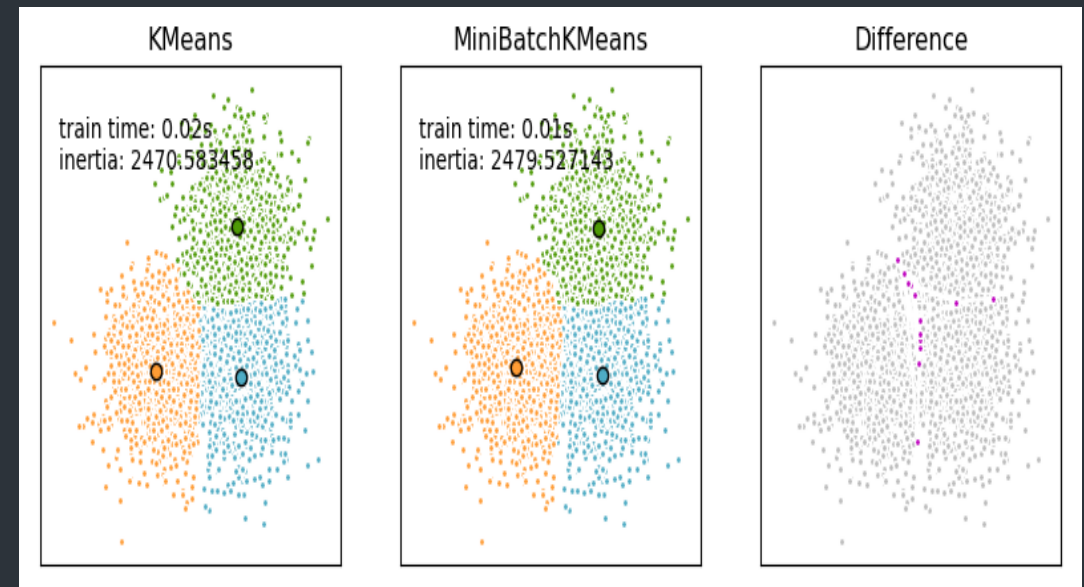
```
from sklearn.cluster import MeanShift
ms = MeanShift(bandwidth=2)
y_pred = ms.fit_predict(X)
```



MiniBatch K-Means

- **How it works:**
 - **Approximate K-Means** using random mini-batches.
 - **Faster but less accurate** than full K-Means.
- **Strengths:**
 - **Scalable to very large datasets.**
 - Useful for **online learning**.
- **Weaknesses:**
 - **Lower accuracy** due to approximations.
- **Best for:**
 - Big data applications (e.g., real-time clustering, recommendation systems).

```
from sklearn.cluster import MiniBatchKMeans  
mbk = MiniBatchKMeans(n_clusters=3, random_state=5)  
y_pred = mbk.fit_predict(X)
```



BIRCH (Balanced Iterative Reducing & Clustering using Hierarchies)

How it works:

- Builds a **Clustering Feature (CF) Tree** for incremental clustering.
- Compresses data into **subclusters** before final clustering.

Strengths:

- Memory-efficient** for large datasets.
- Handles **high-dimensional data** better than K-Means.

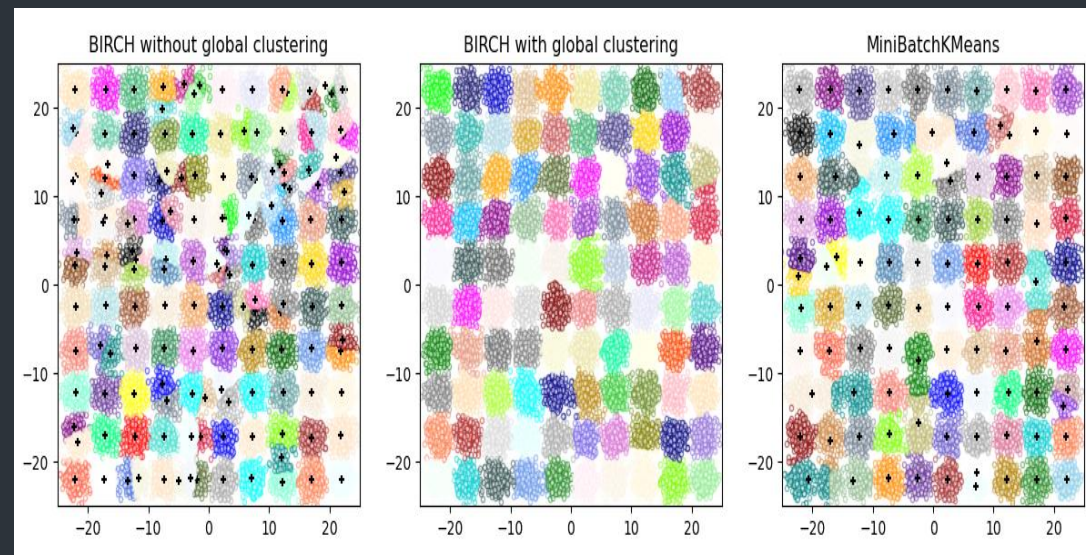
Weaknesses:

- Sensitive to input order** of data.
- Struggles with **non-spherical clusters**.

Best for:

- Large-scale datasets (e.g., customer segmentation, anomaly detection).

```
from sklearn.cluster import Birch
birch = Birch(n_clusters=3)
y_pred = birch.fit_predict(X)
```



HDBSCAN (Hierarchical DBSCAN)

► How it works:

Extends DBSCAN with hierarchical clustering.

Automatically extracts clusters using stability analysis.

► Strengths:

No need for eps parameter (unlike DBSCAN).

Better for varying densities.

► Weaknesses:

Slower than DBSCAN.

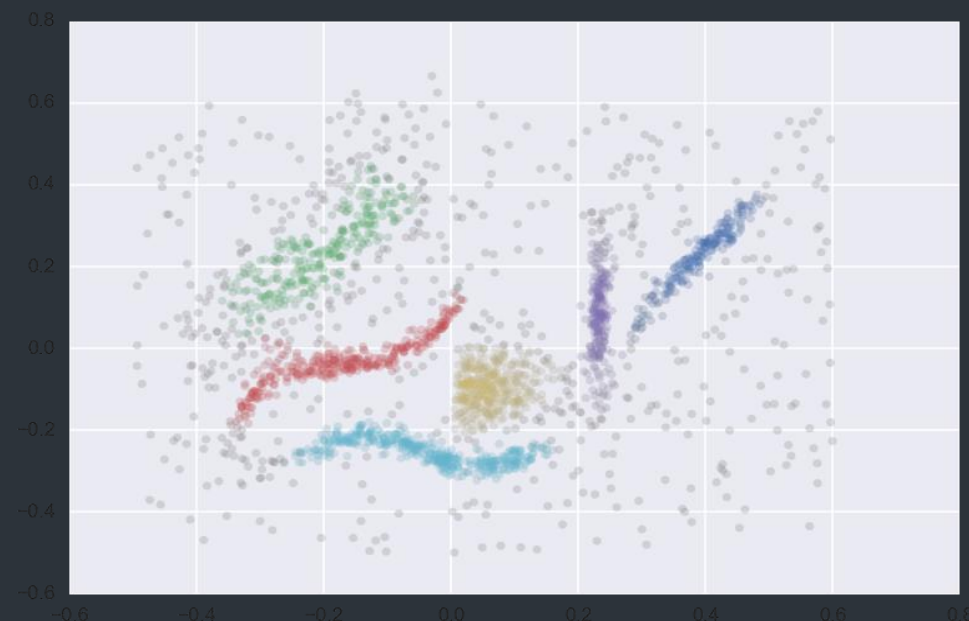
► Best for:

Complex datasets with varying densities (e.g., bioinformatics, social network analysis).

(Requires `hdbscan` library: `pip install hdbscan`)

python

```
import hdbscan
hdb = hdbscan.HDBSCAN(min_cluster_size=5)
y_pred = hdb.fit_predict(X)
```



DBSCAN (Density-Based Spatial Clustering of Applications with Noise)

► How it works:

Groups points based on density (core, border, noise points).

► No predefined cluster count needed.

► Strengths:

Robust to noise and arbitrary cluster shapes.

Works well with spatial data.

► Weaknesses:

Struggles with varying densities.


Sensitive to `eps` and `min_samples` parameters.

► Best for:

Anomaly detection, geographic data, and irregularly shaped clusters.

Rathiga_ML_Cluster

```
from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5)
y_pred = dbscan.fit_predict(X)
```



OPTICS (Ordering Points To Identify Clustering Structure)

► How it works:

Generalized DBSCAN that creates a reachability plot.

Extracts clusters at multiple density levels.

► Strengths:

Handles varying densities better than DBSCAN.

No single eps parameter needed.

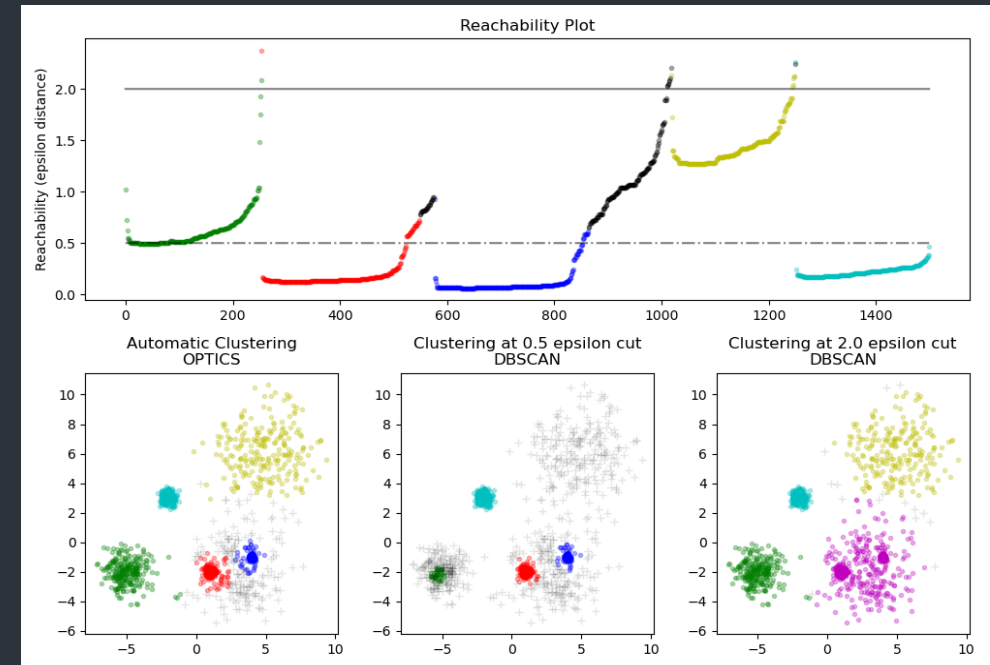
► Weaknesses:

Slower than DBSCAN.

► Best for:

► Datasets with nested clusters (e.g., astronomy, geology).

```
from sklearn.cluster import OPTICS
optics = OPTICS(min_samples=5)
y_pred = optics.fit_predict(X)
```



Spectral Clustering

► How it works:

- Uses **graph Laplacian** to project data into lower dimensions.
- Applies **K-Means on eigenvectors**.

► Strengths:

- Works with **non-convex clusters**.
- Effective for **graph-based data**.


► Weaknesses:

- **Not scalable** ($O(N^3)$ for eigen decomposition).

► Best for:

- Image segmentation, community detection in networks.

```
from sklearn.cluster import SpectralClustering
spec = SpectralClustering(n_clusters=3, random_state=5)
y_pred = spec.fit_predict(X)
```



Summary Table

| Algorithm | Best For | Strengths | Weaknesses |
|---------------------|---------------------------|--------------------------------|--------------------------------|
| AffinityPropagation | Small datasets, unknown K | No need for K, finds exemplars | High computational cost |
| Agglomerative | Hierarchical clustering | Dendrogram, flexible linkage | Slow for large data |
| BIRCH | Large datasets | Memory-efficient, fast | Sensitive to data order |
| DBSCAN | Noise, irregular shapes | No predefined K, robust | Struggles with varying density |
| K-Means | Large, spherical clusters | Fast, scalable | Needs K, sensitive to init |
| HDBSCAN | Varying densities | Automatic K, robust | Slower than DBSCAN |
| Mean Shift | Density peaks | No K needed | Computationally heavy |
| OPTICS | Nested clusters | Multi-density handling | Slower than DBSCAN |
| Spectral | Non-convex clusters | Works on graphs | Not scalable |

Summary Table with Visual Behavior

| Algorithm | Type | Diagram Behavior | Best Use Case |
|----------------|-------------------|-----------------------------------|--------------------------|
| K-Means | Partition-based | Spherical clusters | Well-separated data |
| DBSCAN | Density-based | Arbitrary shapes, noise-resistant | Spatial data, anomalies |
| Agglomerative | Hierarchical | Dendrogram, nested clusters | Taxonomy, bioinformatics |
| Mean Shift | Density-based | Smooth blobs, no fixed K | Image segmentation |
| Affinity Prop. | Exemplar-based | Finds "representative" points | Small datasets |
| Spectral | Graph-based | Non-convex clusters | Image, network data |
| OPTICS | Density-hierarchy | Multi-level density clusters | Nested structures |
| BIRCH | Hierarchical | CF-Tree compression | Large-scale data |

When to Use Which?

- **Default choice: K-Means** (if data is spherical and K is known).
- **Irregular shapes: DBSCAN / HDBSCAN.**
- **Hierarchical structure: Agglomerative / BIRCH.**
- **Unknown K: AffinityPropagation / Mean Shift / HDBSCAN.**
- **Very large data: MiniBatch K-Means / BIRCH.**